

Xenapi Admin Tools 4.0 Scripting Guidelines

(preliminary guide)

All Xenapi Admin Tools 4.0 scripts should include the following functionality

Xenapi Admin Tools 4.0 scripts should

1. display objects by UUID or Name-labels
2. display data in tabulated text format or CSV
3. handle having more than one set of column titles
4. know which directory the tool resides and source library.sh
5. use generic functions from library.sh
6. only loop through UUID lists, never Name-labels
7. hold data in arrays that can be passed to getcolwidth and sort_arrays functions
8. take column padding as a commandline argument
9. display the script name and version in the help
10. include the changelog at the beginning in comments
11. offer an “order by field” option using -o (using sort_arrays function)
12. show UUID/Name pairs sorted by name in select menu (show_arraymenu function)
13. present lists (Name-label | UUID) in alphabetical order (sort_arrays function)
14. be able to present data in color (cecho and setcolors functions)
15. display column padding using printspaces if possible (built in CSV support)
16. operate on remote poolmasters by specifying username/ip or load stored local config (get_poolcreds function)
17. Handle unicode in name-labels e.g. convert unicode to ascii so #{VAR} is correct

Xenapi Admin Tools 4.0 Sections

Below I've used the lsvms.sh script as an example of XCP tools 4.0.

Section 1 – Changelog

```
# Lists XCP/Xenserver Virtual Machines one per line with uuid and host
# Author: Grant McWilliams (grantmcwilliams.com)
# Version: 0.5
# Date: June 27, 2012
# Version: 0.6
# Date: Sept 15, 2012
# Complete rewrite using printspaces, sort_vmnames and getcolwidth
# Now provides four output MODES - name, uuid, mixed and both
# Also provides CSV output
```

Section 2 – Standard functions – setup(), syntax()

The setup() function sets initial variables and executes setcolors(). Setup is always called before getopts in case you want a default variable to be set that may be overridden by a user inputted option on the commandline which is processed by getopts.

```
setup()
{
    SCRIPTDIR=$(dirname $(readlink -f "${BASH_SOURCE[0]}"))
```

```

    source "$SCRIPTDIR/library.sh"
    setcolors
    DEFSPACE="5"
    MINSPE="DEFSPACE"
    MODE="mixed"
    VERSION="0.7"
}

```

The syntax() function gets called from getopt if the option provided by the user is not recognized or is -h. Syntax should display the tool version and its usage.

```

syntax()
{
    echo "$(basename $0) $VERSION"
    echo ""
    echo "Syntax: $(basename $0) [options]"
    echo "Options:"
    echo "-d - shell debugging"
    echo "-c - output comma separated values"
    echo "-u - shows VM UUID, Status, Host UUID"
    echo "-b - shows VM Name, Status, VMUUID, Host Name and Host UUID"
    echo "-n - shows VM Name, Status and Hostname"
    echo "-m - shows VM Name, Status, VM UUID and Hostname"
    echo "-o <value> - changes sort order by column, value can be vmname, hostname, vmuud,      hostuud"
    echo "-s <host> - remote poolmaster host"
    echo "-p <password> - remote poolmaster password"
    echo "-h - this help text"
    echo "-w - number of whitespaces between columns"
    echo ""
    exit
}

```

Section 3 – run setup and get options.

The first item of the main program codebase to run is the setup() function followed by getopt to process the commandline options.

```

setup
while getopt :dcubnmhw:s:o:p: opt ;do
    case $opt in
        d) set -x ;;
        h) syntax ;;
        c) CSV="yes" ;;
        u) MODE="uuid" ;;
        b) MODE="both" ;;
        n) MODE="name" ;;
        m) MODE="mixed" ;;
        o) ORDER="$OPTARG" ;;
        s) REMOTE="yes" ; POOLMASTER="$OPTARG" ;;
        p) PASSWORD="$OPTARG" ;;
        w) isnumber "$OPTARG" && MINSPE="$OPTARG" ;;
        \?) echo "Unknown option"; syntax ;;
    esac
done
shift $(( $OPTIND - 1 ))
getpoolcreds

```

Section 4 – Populating the TITLE[] array

We use an array to store all column titles. We use a case statement to assign different values depending on the commandline options. This is a very flexible system.

```
# Set Title array depending on MODE
case "$MODE" in
    "uuid")  TITLES=( 'VM UUID' 'Status' 'Host UUID' ) ;;
    "name")  TITLES=( 'VM Name' 'Status' 'Host Name' ) ;;
    "mixed") TITLES=( 'VM Name' 'Status' 'VM UUID' 'Host Name' ) ;;
    "both")  TITLES=( 'VM Name' 'Status' 'VM UUID' 'Host Name' 'Host UUID' ) ;;
esac
```

Section 5 – Populating the data arrays

We populate the data arrays in order to pass them to getcolwidth which returns the longest item in the array and in turn populates COLLONGEST[]. It's imperative that the data arrays are synchronized ie. VMNAME[1] is the name of the VM also in VMUUID[1].

I used to loop through the object's UUID to ensure we were dealing with unique items. For instance you could have two VM's named testvm but they will have unique UUID numbers thus looping UUIDs makes sense.

Example of getting data via Xenapi Admin Tools 3.0 (**deprecated**):

```
getvmdata()
VMUUIDS=( $(xe vm-list params=uuid is-control-domain=false --minimal | sed 's/,/\n/g') )
for i in $(seq 0 $((${#VMUUIDS[@]} - 1)) );do
    VMNAMES[$i]=$(xe vm-param-get uuid="${VMUUIDS[$i]}" param-name=name-label)
    STATES[$i]=$(xe vm-param-get uuid="${VMUUIDS[$i]}" param-name=power-state)
    HOSTUUIDS[$i]=$(xe vm-param-get uuid="${VMUUIDS[$i]}" param-name=resident-on)
    if [[ "${HOSTUUIDS[$i]}" = '<not in database>' ]];then
        HOSTUUIDS[$i]="" ; HOSTNAMES[$i]=""
    else
        HOSTNAMES[$i]=$(xe host-param-get uuid="${HOSTUUIDS[$i]}" param-name=name-label)
    fi
done
}
```

This works great but it's very slow because the xe command is being called many times for each iteration. It gets even slower when operating on a remote resource pool. The solution is to create a more comprehensive list and call xe fewer times.

With Xenapi Admin Tools 4.0 we create a more comprehensive list to loop through with multiple fields then use BASH string operators to pull out each item and assign it to the appropriate arrays. There are limitations to when we can do this though. The general rule is that we can create a comprehensive list of multiple items as long as they can be called from ONE xe call. The code is more complex but the speed benefits are immense.

```
getvmdata()
{
```

```

gethostdata
# The last sed replaces every 4th (3 + 1) occurrence of a comma.
# Change the 3 if you have more than 4 columns of data

VMLIST=$(xe vm-list params=uuid,name-label,power-state,resident-on | awk -F: ' '{print $2}' | sed '/^$/d' |
sed -n '1h;2,$H;$ {g;s/\n/,/g;p}' | sed -e's/\([^\,]*\)\{3\}[^\,]*\),/\n/g')
i=0
IFS=$'\n'
for VMLINE in $VMLIST ;do
    VMHOSTUUIDS[$i]="${VMLINE##*,}" ; VMLINE="${VMLINE%,*}"
    STATES[$i]="${VMLINE##*,}" ; VMLINE="${VMLINE%,*}"
    VMNAMES[$i]="${VMLINE##*,}" ; VMLINE="${VMLINE%,*}"
    VMUUUIDS[$i]="${VMLINE##*,}"
    if [[ "${VMHOSTUUIDS[$i]}" = '<not in database>' ]] ;then
        VMHOSTUUIDS[$i]="" ; VMHOSTNAMES[$i]=""
    else
        for j in $(seq 0 $(( ${#HOSTUUIDS[@]} - 1 )) ) ;do
            if [[ "${VMHOSTUUIDS[$i]}" = "${HOSTUUIDS[$j]}" ]] ;then
                VMHOSTNAME[$i]="$HOSTNAMES[$j]"
            fi
        done
    fi
    (( i++ ))
done
}

```

The most important part here is getting the data into a tabulated format so we can use BASH string operators on it. This is also fairly cryptic and hard coded for the output of xe.

```

LIST=$(xe vm-list params=uuid,name-label,power-state,resident-on | awk -F: ' '{print $2}' | sed '/^$/d' | sed -n
'1h;2,$H;$ {g;s/\n/,/g;p}' | sed -e's/\([^\,]*\)\{3\}[^\,]*\),/\n/g')

```

The output of xe vm-list params=uuid,name-label,power-state,resident-on looks like this.

```

uuid ( RO)      : e817b6f8-8896-b453-dead-51bcf5fc5e2a
name-label ( RW): 955275684
power-state ( RO): running
resident-on ( RO): 47c0e22c-eac9-444d-a4b1-d958f0ab60cf

```

The first awk removes the column before the colon (:). The data is then piped to the sed which removes blank lines. The second sed replaces commas with new lines and the **last sed replaces every nth comma with a newline**. Depending on the number of parameters we've specified for xe we need to change where we want sed to break the line. In the case below I'm using the number **3** (4 columns). If you were to have 2 columns you'd change the **3** to a 1.

```

sed -e's/\([^\,]*\)\{3\}[^\,]*\),/\n/g')

```

The idea is that we put all items in CSV then break the CSV file every nth comma so each line has UUID, Name-label, power-state and resident-on. The result is

```

d162eca5-afda-07ba-dc4a-a552673fe8da,987654321,halted,<not in database>
5ac8259d-9240-74a8-f71e-b2118ac88111,987654322,running,47c0e22c-eac9-444d-a4b1-d958f0ab60cf
b6fb3b6f-9281-effb-0e65-9e20b16c1943,CentOS6,halted,<not in database>

```

From this output we can loop over each line and grab the first field using BASH string operators.

```
for LINE in $LIST ;do
    VMUUIDS[$i]="${LINE%%,*}"           ;LINE="${LINE#*,}"
    VMNAMES[$i]="${LINE%%,*}"           ;LINE="${LINE#*,}"
    CONTROLDOMAIN[$i]="${LINE%%,*}"     ;LINE="${LINE#*,}"
done
```

This would split the line at the commas and grab the first field and assign it to \$VMUUIDS. The second half of the line assigns the remaining fields (that we didn't assign) back to \$LINE. Now \$LINE holds the remaining fields and we can use the same method to get each remaining field. BASH string operators can only cut fields from the beginning or end but not the middle so we have to keep assigning the remaining fields back to the \$LINE variable. We could use awk or cut to grab any field we want but echoing the variable and piping it to an external command is much slower than using BASH string operators.

Because of the limitations of xe we have to call it each time we want to create a new list. For instance we couldn't create one list containing output from xe vm-list and output from xe vif-list at the same time easily. So we use this method for EACH unique xe list. It's also important to note that we can usually call xe outside the loop for each list as apposed to calling it multiple times inside of the loop. The time savings can be substantial.

Section 6 – Sort the data arrays

For presentation purposes we want the data arrays sorted alphabetically according to name-label. We use the sort_arrays() function for this.

```
case "$ORDER" in
    vmname) sort_arrays VMNAMES VMUUIDS STATES HOSTUUIDS HOSTNAMES ;;
    hostname) sort_arrays HOSTNAMES VMNAMES VMUUIDS STATES HOSTUUIDS ;;
    vmuuid) sort_arrays VMUUIDS HOSTUUIDS HOSTNAMES VMNAMES STATES ;;
    hostuuid) sort_arrays HOSTUUIDS HOSTNAMES VMNAMES VMUUIDS STATES ;;
esac
```

Section 7 – Get the length of the items in the data arrays

We use a case statement to get the length of each column depending on which commandline option the user chose. This is a very flexible way of switching between using UUIDs or Name-labels for display.

```
# Get the length of each column and store it in COLLONGEST[]
case "$MODE" in
    "uuid") COLLONGEST[0]=$ (getcolwidth "${TITLES[2]}" "${VMUUIDS[@]}")
            COLLONGEST[1]=$ (getcolwidth "${TITLES[1]}" "${STATES[@]}")
            COLLONGEST[2]=$ (getcolwidth "${TITLES[4]}" "${HOSTUUIDS[@]}")
            ;;
    "name") COLLONGEST[0]=$ (getcolwidth "${TITLES[0]}" "${VMNAMES[@]}")
            COLLONGEST[1]=$ (getcolwidth "${TITLES[1]}" "${STATES[@]}")
            COLLONGEST[2]=$ (getcolwidth "${TITLES[3]}" "${HOSTNAMES[@]}")
            ;;
    "mixed") COLLONGEST[0]=$ (getcolwidth "${TITLES[0]}" "${VMNAMES[@]}")
            COLLONGEST[1]=$ (getcolwidth "${TITLES[1]}" "${STATES[@]}")
            COLLONGEST[2]=$ (getcolwidth "${TITLES[2]}" "${VMUUIDS[@]}")
```

```

COLLONGEST[3]=$ (getcolwidth "${TITLES[3]}" "${HOSTNAMES[@]}")
;;
"both") COLLONGEST[0]=$ (getcolwidth "${TITLES[0]}" "${VMNAMES[@]}")
COLLONGEST[1]=$ (getcolwidth "${TITLES[1]}" "${STATES[@]}")
COLLONGEST[2]=$ (getcolwidth "${TITLES[2]}" "${VMUUIDS[@]}")
COLLONGEST[3]=$ (getcolwidth "${TITLES[3]}" "${HOSTNAMES[@]}")
COLLONGEST[4]=$ (getcolwidth "${TITLES[4]}" "${HOSTUUIDS[@]}")
;;
esac

```

Section 8 – Printing the column headings and data columns

The last section outputs the column headings and data columns to the user. We loop through the VM UUIDS (never loop through name-labels) and using a case statement to print out data differently depending on the options specified by the user. We use `cecho` to print in color and `printspaces` to print out the padding between columns. Because we populated the arrays in section 7 we don't have to do anything but present data here. This is not always 100% possible though. On occasion you may also have to gather data in this loop.

```

# Print column headings and data columns
printheadings
for i in $(seq 0 $(( ${#VMUUIDS[@]} - 1 )) );do
    case "$MODE" in
        "uuid")
            cecho "${VMUUIDS[$i]}" cyan
            printspaces "${COLLONGEST[0]}" "${#VMUUIDS[$i]}"
            cecho "${STATES[$i]}" red
            printspaces "${COLLONGEST[1]}" "${#STATES[$i]}"
            cecho "${HOSTUUIDS[$i]}" blue
            ;;
        "name")
            cecho "${VMNAMES[$i]}" cyan
            printspaces "${COLLONGEST[0]}" "${#VMNAMES[$i]}"
            cecho "${STATES[$i]}" red
            printspaces "${COLLONGEST[1]}" "${#STATES[$i]}"
            cecho "${HOSTNAMES[$i]}" blue
            ;;
        "mixed")
            cecho "${VMNAMES[$i]}" cyan
            printspaces "${COLLONGEST[0]}" "${#VMNAMES[$i]}"
            cecho "${STATES[$i]}" red
            printspaces "${COLLONGEST[1]}" "${#STATES[$i]}"
            cecho "${VMUUIDS[$i]}" blue
            printspaces "${COLLONGEST[2]}" "${#VMUUIDS[$i]}"
            cecho "${HOSTNAMES[$i]}" blue
            ;;
        "both")
            cecho "${VMNAMES[$i]}" cyan
            printspaces "${COLLONGEST[0]}" "${#VMNAMES[$i]}"
            cecho "${STATES[$i]}" red
            printspaces "${COLLONGEST[1]}" "${#STATES[$i]}"
            cecho "${VMUUIDS[$i]}" blue
            printspaces "${COLLONGEST[2]}" "${#VMUUIDS[$i]}"
            cecho "${HOSTNAMES[$i]}" blue
            printspaces "${COLLONGEST[3]}" "${#HOSTNAMES[$i]}"

```

```

                cecho "${HOSTUUIDS[$i]}" blue
            ;;
        esac
        echo ""
    done

```

Xenapi Admin Tools 4.0 Functions in Depth

setcolors() and cecho()

The `setcolors()` function assigns ansi values to variables. The `cecho()` function below uses these variables to echo messages in color. To set a color use `cecho "Hello World in Blue" blue`. To turn off the ansi values use the off keyword eg. `cecho "Hello world" off`. The `cecho` function will display a quoted message in the color specified. It does not print a carriage return so must be followed by an `echo ""` if one is needed.

```

setcolors
cecho "Warning!!" red

```

getcolwidth()

The `getcolwidth()` function takes an array as an argument and returns an integer representing the length of the longest item line the array. For instance if you had an array that had these items

```

array[0]=bob
array[1]=sally
array[2]=rumplestilskin

```

and you wanted to know how long the longest item is (`array[2]` you'd pass the entire thing to `getcolwidth` like this.

```

COLLONGEST[0]=$(getcolwidth "${TITLES[2]}" "${VMUUIDS[@]}")

```

The `getcolwidth()` function calculates the longest item in `TITLES[2]` and `VMUUIDS[@]` and returns an integer which is assigned to `COLLONGEST[0]`. This way I have the longest item in column one in `COLLONGEST[0]`, the longest item in column two in `COLLONGEST[1]` and so on.

printspaces()

The `printspaces()` function takes two arguments – the longest item in a column returned from `getcolwidth()` and the length of the current item. To check the length of `$VAR` use the following syntax `${#VAR}`. Below we're echoing `${VMNAME[1]}`, then passing `${COLLONGEST[1]}` and `${#VMNAME[1]}` to `printspaces()`. `Printspaces()` will subtract the length of `${VMNAME[1]}` from `${COLLONGEST[1]}` and print that amount of empty spaces for proper column padding.

```

cecho "${VMNAME[1]}" red ; printspaces "${COLLONGEST[1]}" "${#VMNAME[1]}"

```

printheadings()

The printheadings() function prints out the column headings previously assigned in TITLES[]. Just call it with no arguments.

```
printheadings
```

sort_arrays()

The sort_arrays() function sorts multiple arrays in parallel. The first array passed to it is used to determine the order of all array items. Whatever is done to the first array will be done to the others. For example you have two arrays indexed together.

```
VMUUIDS[1]="8ea244fd-b67a-78c2-9317-18adb29ae8cc"  
VMUUIDS[2]="4303c7e3-0479-42bc-c115-5a3877f1403f"  
  
VMNAMES[1]="webserver"  
VMNAMES[2]="mailserver"
```

In this case the UUID for “webserver” is 8ea244fd-b67a-78c2-9317-18adb29ae8cc because the array index for both is 1 (VMUUIDS[1] and VMNAMES[1]). We have to have multiple arrays indexed together with Bash 2 and 3 because they don't support associative arrays. To sort both of these arrays and keep the associations correct we'd use the sort_arrays() function. Note: you specify the array NAME only eg. VMNAMES not \$VMNAMES or VMNAMES[@].

```
sort_arrays VMNAMES VMUUIDS
```

fsort_arrays()

The fsort_arrays() function works just like sort_arrays and even uses the same syntax. The difference is that it prints out the value of the arrays passed into a temporary file and then uses the sort command to sort it.

```
fsort_arrays VMNAMES VMUUIDS
```

yesno()

The yesno() function just asks a simple question and returns a 0 or 1 depending on the answer. yesno() takes one argument – the message to display.

To use yesno() you just place it in an if conditional followed by the message to display.

```
If yesno "Do you like dogs?" ; then  
    echo "you like dogs"  
else  
    echo "You don't like dogs"  
fi
```


isnumber()

The isnumber() function checks to see if the value in the variable only contains an integer. If any other characters are found it will return 1.

```
If isnumber "43" ; then
    echo "is a number"
else
    echo "is not a number"
fi
```

getunit()

The getunit() function will convert an integer of bytes to Kilobytes, Megabytes, Gigabytes or Terabytes. It returns the number followed by an abbreviation ie. 964M or 123G.

To use it just pass the integer to the getunit() function.

```
SIZE=$(getunit 9456)
```

getpoolcreds()

The getpoolcreds() function checks for the existence of \$SCRIPTDIR/.XECONFIG". The .XECONFIG directory holds one config file per pool. Inside each config file are lines for poolmaster, username, port and password.

```
POOLMASTER="cloud0"
PORT="443"
USERNAME="root"
PASSWORD="password"
```

The permissions are checked on the .XECONFIG directory and if they're readable by anyone else besides root the script throws an error and exits. This behavior may change in the future and offer the user the chance of modifying the permissions.

If .XECONFIG exists getpoolcreds() parses each text file whose name ends in .cfg for hostname, port, username and password. If any items are not set it prompts the user for the value. The variable XE_EXTRA_ARGS is set and exported globally. If no configs are found it assumes the poolmaster is localhost.

Using getpoolcreds is as simple as just calling it with no arguments.

```
getpoolcreds
```

show_arraymenu()

The show_arraymenu() function is a semi-hard coded function that takes at least one array as an

argument and presents a select menu to the user. When the user selects an item `show_arraymenu()` strips any valid UUID from that selection and returns it.

The purpose of such a function would be to present the user with a menu of Name-label and UUID pairs, when the user selects a Name-label (UUID is for when there are multiple identical name-labels) the UUID is returned. It allows a very easy way to present the user with a sorted list of name-labels and then return a reliable UUID for further processing.

Pass at least one array when calling `show_arraymenu()`

```
show_arraymenu VMNAMES VMUUIDS
```

Note 1: only pass the array names not the arrays themselves ie. `VMNAMES` not `$VMNAMES` or `{VMNAMES[@]}`.

Note 2: do NOT pass more than one UUID array to `show_arraymenu` as it will attempt to strip ANY UUIDs and return them. It only handles one UUID.

gethostdata()

Get host list and assign `HOSTUUIDS[]` and `HOSTNAMES[]` arrays. Call it directly to populate the `HOSTUUIDS` and `HOSTNAMES` arrays. They will be synchronized so `HOSTUUIDS[1]` is the UUID of `HOSTNAMES[1]`. The arrays will be alphabetically sorted based on the values of `$HOSTNAMES`.

```
gethostdata
for i in ${HOSTUUIDS[@]} ;do
    echo "$i"
done
```

getcpudata()

```
# Get CPU core list and assign CPUUUIDS[] and CPUNUMS[] arrays
```

Wishlist for Xenapi Admin Tools 5.0

1. general `getdata` function to replace all the `getsrdata`, `getvmdata` etc. Syntax would be `getdata "xe commands to execute"`. The results would be placed in arrays with the same names as the paramaters e.g. `xe vm-list params=uuid,name-label` would set the values of `UUID[1]` and `NAME-LABEL[1]`. Obviously there's a namespace issue so whenever using `getdata` you'd need to utilize all of the results ie. echoing them or reassigning them before calling it again. Below is the general idea.

```
getdata()
{
```

```

COMMAND="$1"
NUMARGS=$(echo $COMMAND | awk -F'params=' '{print $2}' | <do something to count fields>
local LIST=$(($COMMAND | awk -F': ' '{print $2}' | sed '/^$/d' | sed -n '1h;2,$H;${g;s/\n/,/g;p}' | sed -e's/\
([^\,]*,)\{ $NUMARGS\} [^\,]*\),/1\n/g')
i=0
for LINE in $LIST ;do
    <do something to get name of arg>
    eval ARGNAME[$i]="${LINE%%,*}" ;LINE="${LINE#*,}"
    (( i++ ))
done
sort_arrays HOSTNAMES HOSTUUIDS
}

```